



# First Order Logic with Logic Tensor Networks

---

Federico Bianchi  
**Bocconi University**  
[f.bianchi@unibocconi.it](mailto:f.bianchi@unibocconi.it)



# Logic Tensor Networks

---

# Logic Tensor Networks

## Logic Tensor Networks: Deep Learning and Logical Reasoning from Data and Knowledge<sup>★</sup>

Luciano Serafini<sup>1</sup> and Artur d'Avila Garcez<sup>2</sup>

<sup>1</sup> Fondazione Bruno Kessler, Trento, Italy, [serafini@fbk.eu](mailto:serafini@fbk.eu)

<sup>2</sup> City University London, UK, [a.garcez@city.ac.uk](mailto:a.garcez@city.ac.uk)

# Logic Tensor Networks

Logic Tensor Networks [Serafini+,2016] (LTNs)

**LTNs = Neural Networks + First Order Fuzzy Logic**

**Key Aspects:**

- LTNs **ground fuzzy logic in a vector space: continuous values in [0,1]**
- LTNs assign truth values to formulas using neural networks
- LTNs can learn from both data and rules
- LTNs can be used to do inferences over rules after training

**Key Idea:** LTNs provide a method to learn reasoning over vector spaces

# Logic Tensor Networks

Logic Tensor Networks [Serafini+,2016] (LTNs)

**LTNs = Neural Networks + First Order Fuzzy Logic**

**Key Aspects:**

- LTNs ground fuzzy logic in a vector space: continuous values in [0,1]
- LTNs **assign truth values to formulas using neural networks**
- LTNs can learn from both data and rules
- LTNs can be used to do inferences over rules after training

**Key Idea:** LTNs provide a method to learn reasoning over vector spaces

# Logic Tensor Networks

Logic Tensor Networks [Serafini+,2016] (LTNs)

**LTNs = Neural Networks + First Order Fuzzy Logic**

**Key Aspects:**

- LTNs ground fuzzy logic in a vector space: continuous values in [0,1]
- LTNs assign truth values to formulas using neural networks
- LTNs **can learn from both data and rules**
- LTNs can be used to do inferences over rules after training

**Key Idea:** LTNs provide a method to learn reasoning over vector spaces

# Logic Tensor Networks

Logic Tensor Networks [Serafini+,2016] (LTNs)

**LTNs = Neural Networks + First Order Fuzzy Logic**

## Key Aspects:

- LTNs ground fuzzy logic in a vector space: continuous values in [0,1]
- LTNs assign truth values to formulas using neural networks
- LTNs can learn from both data and rules
- LTNs **can be used to do inferences over rules after training**

**Key Idea:** LTNs provide a method to learn reasoning over vector spaces

# Logic Tensor Networks: Plan

We need:

- A way to put logic in the vector space;
- A way to learn logical representation;
- A way to compute the truth score of predicate.

What we get:

- **A recursive logical language that works in the vector space.**

# Logics in the Vector Space


---

# Logic Tensor Networks: General Idea

**parent(Susan, Ann)**

# Logic Tensor Networks: General Idea


parent(Susan, Ann)



Constants are  
points in  $\mathbb{R}^k$

# Logic Tensor Networks: General Idea


parent(Susan, Ann)




These are logical constants, but **we can learn these** vectors (similarly to what we do with word embeddings) or **initialize them** with pre-trained embeddings

Constants are points in  $\mathbb{R}^k$


# Logic Tensor Networks: General Idea



# Logic Tensor Networks: General Idea



# Logic Tensor Networks: General Idea




# Logic Tensor Networks: General Idea

**parent(Susan, Ann) & parent(Mike, Robert)**




# Logic Tensor Networks: General Idea

**parent(Susan, Ann) & parent(Mike, Robert)**




# Logic Tensor Networks: General Idea

$\text{parent}(\text{Susan}, \text{Ann}) \& \text{parent}(\text{Mike}, \text{Robert})$



# Logic Tensor Networks: General Idea

$\text{parent}(\text{Susan}, \text{Ann}) \& \text{parent}(\text{Mike}, \text{Robert})$




# Learning Logics in the Vector Space

---

# Example

KB:

- !parent(mark, john)
- parent(john,mark)
- ancestor(mark, lucas)
- parent(john, susan) | parent(john, dania)



KB:

- !parent(mark, john)
- parent(john, mark)
- ancestor(mark, lucas)
- parent(john, susan) | parent(john, dania)


# Example: Forward Pass

1 -



$$\begin{array}{c|c} \text{M} & \text{J} \\ \hline \text{---} & \text{---} \\ \text{---} & \text{---} \end{array} = 0.8$$

We want this axiom to be true, we need to maximize this




KB:

- !parent(mark, john)
- parent(john, mark)
- ancestor(mark, lucas)
- parent(john, susan) | parent(john, dania)

# Example: Back Pass

1 -




$$\begin{array}{c|c} \text{M} & \text{J} \\ \hline \text{---} & \text{---} \\ \hline \end{array} = 0.8$$


A diagram showing two vertical columns of five green squares each. Below the first column is the letter "M" and below the second is the letter "J". To the right of the columns is the equation " $= 0.8$ ".

Error is 0.2 (with MAE)

Update using backpropagation




# Example: Forward and Back Pass



We want to maximize this and thus we update the respective values

KB:

- !parent(mark, john)
- **parent(john,mark)**
- ancestor(mark, lucas)
- parent(john, susan) | parent(john, dania)



KB:


- !parent(mark, john)
- parent(john, mark)
- ancestor(mark, lucas)
- **parent(john, susan) | parent(john, dania)**

# Example: Forward and Back Pass



$$\max(0.2, 0.9) = 0.9$$

We want to maximize this and thus we update the respective values



# A few more things

---

# Logic Tensor Networks: Quantified Formulas

How do we interpret a quantified formula?

- $\forall x \text{ human}(x)$

In LTN each variable ( $x$ ) is associated to a **domain sample**

- $S = \{\text{John}, \text{Paul}\}$  and  $x$  is defined over  $S$ , truth value of  $\forall x \text{ human}(x)$  is
  - **min human( $x$ ),  $\forall x \in S$ .**
    - $\text{human}(\text{John}) = 0.8, \text{human}(\text{Paul}) = 0.9, \forall x \text{ human}(x) = 0.8$
  - the truth value of the forall of a predicate is the value of the least true axiom.

# Logic Tensor Networks: Data and Rules

LTNs can learn from both data and rules.

- Quantifiers are defined **over a domain sample**.


**parent(Mark,Susan)**     $\forall x,y \text{ parent}(x,y) \rightarrow$   
**parent(Ron,Susan)**    **ancestor(x,y)**

Quantifiers interpreted using an **aggregation function** (e.g., average):  
 $\forall x P(x)$  = average value of  $P(x)$  in LTNs.

# Logic Tensor Networks: Learning

The network is trained on a **best satisfiability task**:

- Learn the representations of:
  - **vectors** for the constants, **parameters** for the predicates in such a way that **the axioms are satisfied in the best possible way**.



Given **parent(Ann, Susan)** we expect the network to **learn representations** for **Ann, Susan and parent** in such a way that the predicted value is close to 1

## Logic Tensor Networks: After Training Inference

- The trained network can be used to make novel inferences.
- Suppose we train using a dataset of ***parents*** and ***ancestors*** relationships.

# Logic Tensor Networks: After Training Inference

- The trained network can be used to make novel inferences.
- Suppose we train using a dataset of ***parents*** and ***ancestors*** relationships.
- **After training** we can query LTNs on:
  - $\forall x,y \text{ ancestor}(x,y) \rightarrow \text{parent}(x,y)$  has truth value close to 0

We get a deep recursive logical language to explore the vector space!


# How to Compute The Truth Score

---

# Logic Tensor Networks: Neural Tensor Network

$$g(e_1, R, e_2) = u_R^T f \left( e_1^T W_R^{[1:k]} e_2 + V_R \begin{bmatrix} e_1 \\ e_2 \end{bmatrix} + b_R \right)$$

Compute truth score  
 $R(e_1, e_2)$  



Socher, R., Chen, D., Manning, C. D., & Ng, A. (2013). Reasoning with neural tensor networks for knowledge base completion. In Advances in neural information processing systems (pp. 926-934).

# Logic Tensor Networks: LTN's Neural Tensor Network

$$\mathcal{G}(P) = \sigma \left( u_P^T \tanh \left( \mathbf{v}^T W_P^{[1:k]} \mathbf{v} + V_P \mathbf{v} + B_P \right) \right)$$

- P is the axiom (e.g., **parent(Ann, Susan)**);
- **v** is **the concatenation** of the vectors of the arguments of the predicate.
  - Concatenation is order dependent =>  $\text{parent(Ann,Susan)} \neq \text{parent(Susan,Ann)}$


# Logic Tensor Networks: LTN's Neural Tensor Network

$$\mathcal{G}(P) = \sigma \left( u_P^T \tanh \left( \mathbf{v}^T W_P^{[1:k]} \mathbf{v} + V_P \mathbf{v} + B_P \right) \right)$$

- P is the axiom (e.g., **parent(Ann, Susan)**);
- **v** is **the concatenation** of the vectors of the arguments of the predicate.
  - Concatenation can be used for higher arity-predicates like
    - **married(John, Susan, Saint Paul Church, 10/10/1984)**


# Logic Tensor Networks: LTN's Neural Tensor Network

- married(John, Susan, Saint Paul Church, 10/10/1984)



# Logic Tensor Networks: LTN's Neural Tensor Network


- married(John, Susan, Saint Paul Church, 10/10/1984)



# Logic Tensor Networks: Complex Formulas

- $(\text{parent}(\text{John}, \text{Susan}) \& \text{parent}(\text{Susan}, \text{Tom})) \Rightarrow \text{parent}(\text{John}, \text{Tom})$

Complex formulas are interpreted **recursively**.



# Recap

- **LTNs** grounds logic in the vector space: each predicate is an NN that we train.
- **LTNs** gives us a recursive logical language
- We can **logically explore** the space after training with compositions of learned predicates.

# Logic Tensor Networks with Embeddings

---

# Support Logical Reasoning with Embeddings

Agent axiomatic knowledge:

```
species(cat)  
mammal(tiger)  
bird(penguin)  
 $\forall x (mammal(x) \rightarrow animal(x))$ 
```

logically Inferring something about  
**cat** is not possible

Agent “distributional knowledge”: **cats** and **tigers** are more similar than **cats** and **penguins**.



Possible inferences with the support of distributional knowledge:

**mammal(cat)**

But also:

**animal(cat)**

**Key Idea:** combining standard logical reasoning with embeddings

# Sub-symbolic Commonsense: Distributional Semantics

**Distributional Hypothesis:**  
**similar** words tend to appear in similar **contexts**

[Wittgenstein1953, Firth1957]


What is the meaning of **bardiwac?** [Lenci&Evert]:

- The drinks were delicious: blood-red **bardiwac** as well as light, sweet Rhenish
- He handed her a glass of **bardiwac**
- Beef dishes are made to complement the **bardiwacs**

**Key Idea:** meaning can be derived from language usage

# Sub-symbolic Commonsense: Word2Vec

- Grounded in Distributional Hypothesis [Harris,1954]: vector representations of language items
- i.e., embeddings, generated from a text corpus [Mikolov+, 2013]



- Neural network trained on a prediction task
- Similar words appear in similar contexts and have similar vectors

**Key Idea:** neural models generate distributional vectors of a given dimension

# **Logical Reasoning with Common Sense:**

## Combining LTNs with Embeddings

---

# LTNs and Commonsense

## LTNs

- Learn from structured knowledge
- **Logic constants are vectors**
- Reasoning done combining fuzzy logics with neural networks


## Knowledge Embeddings

- Learned from text
- Capture latent semantics
- **Entities are vectors**


We can use sub-symbolic representations as the constants in LTNs

# Logic Tensor Networks: General Idea

No need to learn the representation of “dbr:cat”. We already have its **distributional entity vector**.



# LTNs and Commonsense




# Experimental Evaluation

We created a simple KB with predicates about **species** and **sub-classes**.


## Task: KB completion.

Inputs to LTNs:

- **Entity Embeddings** from Wikipedia
- **Axiomatic Knowledge (both inst. atoms and 22 quantified rules)**




General Idea




# Experiments: Standard Training/Testing

F1 Scores



# Experiments: Zero-Shot on Novel Entities

F1 Scores



D3

# LTN: Testing New Predicates

We collected species and trained LTNs to recognize the classes of each species

Again: we can **logically explore** the vector space

We query the vector space with logical rules.

Axiom	Truth
$\forall x(species(x) \rightarrow animal(x))$	0
$\forall x(eukaryote(x) \rightarrow \neg bacteria(x))$	0.73
$\exists x(eukaryote(x) \wedge \neg plant(x))$	1

# LTN: Testing New Predicates

We collected species and trained LTNs to recognize the classes of each species

Again: we can **logically explore** the vector space

We query the vector space with logical rules.

Axiom	Truth
$\forall x(species(x) \rightarrow animal(x))$	0
$\forall x(eukaryote(x) \rightarrow \neg bacteria(x))$	0.73
$\exists x(eukaryote(x) \wedge \neg plant(x))$	1

# LTN: Testing New Predicates

We collected cities and people from Wikipedia and learned some predicates like **bornIn**, **locatedIn** and **nationality**

Again: we can **logically explore** the vector space

Axiom	Truth
$\forall x, y, z(nationality(x, y) \wedge locatedIn(z, y) \rightarrow bornIn(x, z))$	0.33
$\exists x(nationality(x, Canada) \wedge bornIn(x, Montreal))$	1
$\forall x(bornIn(x, New\ York) \rightarrow nationality(x, United\ States))$	0.88

# Thank you!

---

# Backup Slides

---

# Terminological Recap: First Order Logic

A **constant** is an element of a domain (set) taken in consideration

$$\begin{aligned}S &: \{John, Paul, \dots\} \\T &: \{BMW, Toyota, \dots\}\end{aligned}$$

A **function** is a relation  $f: S \rightarrow T$  between sets that associates to every element of a first set exactly one element of the second set.

$$\begin{aligned}car: T &\rightarrow S \\car(Paul) &= BMW\end{aligned}$$

A **predicate** is a Boolean-valued function  $P: S \rightarrow \{1 (= \text{True}), 0 (= \text{False})\}$ .

$$\begin{aligned}male: S &\rightarrow \{1, 0\} \\male(John) &= 1\end{aligned}$$

$$\begin{aligned}drives: S \times T &\rightarrow \{1, 0\} \\drives(Paul, BMW) &= 1\end{aligned}$$

# Terminological Recap: First Order Logic

An **axiom** is a statement in a logical language:

$$R(a, b)$$

A **grounded axiom** contains grounded constants:

$$\text{parentOf(John, Paul)}$$

A **quantified axiom** is an axiom that contains quantified variables:

$$\forall x \text{Human}(x)$$

A **formula** is a combination of grounded and quantified axioms:

$$\forall x,y \text{parentOf}(x, y) \Rightarrow \text{parentOf}(y, x)$$

# Terminological Recap: Operators

- & - conjunction:  $\text{parent}(\text{Ann}, \text{Susan}) \ \& \ \text{Female}(\text{Susan})$
- | - disjunction:  $\text{parent}(\text{Ann}, \text{Susan}) \ | \ \text{parent}(\text{Susan}, \text{Ann})$
- ! - negation:  $!\text{Male}(\text{Susan})$
- $\Rightarrow$  - implication:  $\text{parent}(\text{Ann}, \text{Susan}) \Rightarrow \text{children}(\text{Susan}, \text{Ann})$
- $\forall$  - forall:  $\forall x \text{ female}(x) \Rightarrow \text{human}(x)$  [*forall x if x is a female than x is a human*]
- $\exists$  - exists:  $\exists x \text{ human}(x)$  [*exists one human*]